

Byzantine Fault-Tolerant and Locality-Aware Scheduling MapReduce

Andrea Graziani

andrea.graziani93@outlook.it

Università Degli Studi di Roma Tor Vergata

Rome, Italy

ABSTRACT

MapReduce is a programming paradigm developed by Google that enables massive scalability across hundreds or thousands of servers allowing to process very large data-set [3].

Both academic literature and daily experience shows that *arbitrary faults* frequently occur corrupting our results [1]. Moreover, ignore data-set locality can lead to performance degradation and a pointless bigger network traffic [2].

In this paper we present an original MapReduce runtime system capable both to tolerate arbitrary faults and to recognize input data network locations and sizes in order to perform a *locality-aware task scheduling*, improving performance and diminishing network traffic.

Although job's execution using our system requires more resources respect to other implementations, we believe that this cost is acceptable for critical applications that require an higher level of fault tolerance.

KEYWORDS

MapReduce, Fault tolerance, Arbitrary failure, Data locality

1 INTRODUCTION

Various data-intensive tasks, like seismic simulation, natural language processing, machine learning, astronomical data parsing, web data mining and many other, require a processing power that exceeds the capabilities of individual computers imposing the use of a *distributed computing* [2].

Nowadays many famous distributed applications use MapReduce framework as solution for processing large data sets in a distributed environment. In order to properly provide services to an increasing number of users worldwide, is necessary to connect together thousands of computers and hundreds of other devices like network switches, routers and power units, moving consequently an huge amount of data between computers.

However, as many studies confirm, *hardware component failures are frequent* and they will probably happen more often in the future due to increasing number of computers connected to internet [1]. Is been documented that in the first year of a cluster at Google there were 1000 individual machine failures and thousands of hard drive failures [5]. A recent study of DRAM errors in a large number of servers in

Google data-centres for 2.5 years concluded that these errors are more prevalent than previously believed, with more than 8% DIMM affected by errors yearly, even if protected by error correcting codes (ECC) [6]. A Microsoft study of 1 million consumer PCs shown that CPU and core chipset faults are also frequent. [4] Moreover moving large amount of data repeatedly to distant nodes is becoming the bottleneck owing to an increased network traffic causing performance degradation.

Therefore to construct a distributed system capable to provide its services even in the presence of failures is become so critical; consequently, to provide a *fault tolerant* cloud application represents an important goal in distributed-systems design.

Moreover, due to very large data-set involved during data-intensive tasks, exploiting data locality, in order to mitigate network traffic and delay, becomes very important to improve performance.

In these paper we will describe our *Arbitrary Fault-Tolerant Locality-Aware* (AFTLA) MapReduce system capable to resolve all issues described above. We have implemented an AFTLA MapReduce system, capable to perform a word-count of a given text input, from scratch using Go programming language, adopting Apache ZooKeeper™ for maintaining configuration information and system coordination¹; all libraries used in our implementation are fully described in table 1. Our prototype's source code is fully available on GitHub², famous web-based hosting service for version control using `git`; we remind that \LaTeX source code of this paper is available too³.

2 ASSUMPTIONS

In order to properly describe how our algorithm can tolerate arbitrary faults and to exploit data locality, it is necessary to describe its architecture. To do that, in this section we start making some very important assumptions.

Our system is composed by a *set of distributed processes*, every of which run on different hosts in a same data-center; in our prototype, every process run on his own Amazon EC2

¹See <https://zookeeper.apache.org>

²See <https://github.com/AndreaG93/SDCC-Project>

³See <https://github.com/AndreaG93/SDCC-Project-Report>

Table 1: Libraries used in our implementation

Name	Description	Link
go-zookeeper/go	Native Go Zookeeper Client Library	https://github.com/samuel/go-zookeeper
logrus	A structured logger for Go	https://github.com/Sirupsen/logrus
aws-sdk-go	AWS SDK for the Go programming language.	https://github.com/aws/aws-sdk-go

server hosted in a same region (us-east-1). All processes interact with each other using RPC

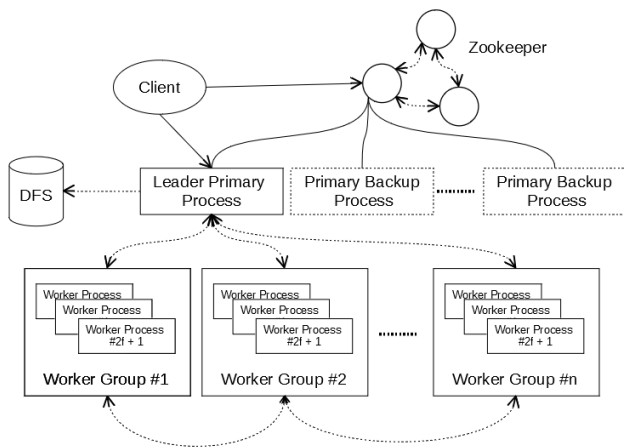
We assume that our system runs *asynchronously*, that is no assumptions about process execution speeds or message delivery times are made; therefore we can normally use time-outs to state that a process has crashed although, occasionally, such conclusion is false [7].

All processes are connected by *reliable channels*, so no messages are lost, duplicated or corrupted; that feature is guaranteed by TCP connections. All system processes use RPC protocol in order to interact with each other.

Clients are *always correct*, because if they are not there is no point in worrying about the correctness of our system's output. Finally we assume the existence of a hash function that is *collisions-resistant*, for which it is infeasible to find two inputs that produce the same output.

In this paper, we take for granted reader's knowledge of the MapReduce framework and Hadoop implementation framework.

3 SYSTEM'S MODEL

**Figure 1: System Architecture**

Three kind of process are used in our system:

Client Processes Which send word-count requests to our system sending directly an input text file.

Worker Process (WP) They execute *map* or *reduce* scheduled by LPP.

Primary Process (PP) It is similar to *JobTracker* process used in Apache Hadoop which duty is to satisfy clients requests scheduling *map* or *reduce* tasks.

Primary Process

At this point, is necessary to make some clarifications about PP's features.

Unlike Hadoop, according to which *JobTracker* process is assumed always correct, the host where PP is running may fail, for example by crashing or by losing network connectivity; in other words, as known in academic literature, PP represents a *single point of failure*.

Therefore, to ensure an high system availability, adopting the same design used in Apache MesosTM ⁴, our system architecture is based on multiple PP backup copies run on different machines; using any leader election algorithm, is possible to elect one of these copies as *Leader Primary Process* (LPP) while all the others remain on stand-by. Using this approach, is possible to manage PP crash failure: when current leader fails, a backup copy is promoted to become the new .

LPP represents our system coordinator since only it can interact with WP scheduling task using a *push-based* approach (unlike Hadoop according to which *JobTracker* schedules task using a *pull-based* approach [2])

However, to perform its task, LPP manages various information about clients requests, many of which, as we will see later, are sent by WP after a map or reduce task. In other word, LPP has got a status.

Is very important to specify that LPP status is stored in memory, therefore they are permanently lost after a crash; this design makes our system easier to implement and helps to reducing overhead due to the disk I/O activities. However recover lost in-memory leader state after a failure is required in order to satisfy clients requests.

Therefore we have adopt a design according to which LPP make a sort of backup, or snapshot, of its current state regularly. To make our implementation easier, in our prototype, considering that LPP state contains a very small amount of data, LPP stores its state in a Zookeeper cluster after some events, like a map task termination. In case of crash failure of current LPP, when a primary process backup becomes

⁴See <http://mesos.apache.org/documentation/latest/architecture/>

leader, it retrieves all data from Apache ZooKeeper cluster, recovering last saved state.

Worker Process and Groups

To make WP arbitrary fault-tolerant we have adopt a design according to which each task (map, reduce, etc.) is executed more than once by different processes running on different host. This approach requires to organize all our WPs into several *groups* through which became possible to mask one or more faulty processes of a given group.

Before explaining how we manage WP faults, we must explain what is a Worker Processes Group (WPG). A WPG is a set of identical worker processes, each of which execute, relatively to a given job, *the same commands using the same input data in the same order*; therefore we expected that all WPs belonging to a same group would produce the same output when no arbitrary fault occurs. All commands and data are sent by LPP.

Suppose to want tolerate at most f arbitrary failures of all WPs belonging to a given group.

To achieve that level of fault-tolerance, we can apply to reach, for instance, *byzantine agreement* among all processes belonging to group. However we believe that solution is too expensive because it requires, as properly explain in literature[7], $3f + 1$ replicas.

Therefore we have adopt a protocol based on *majority voting* inspired to quorum-based protocols applied in replicated-write systems. As explained in literature[7], in this kind of systems, if processes exhibit arbitrary failures, continuing to run when faulty and sending out erroneous or random replies, a minimum of $2f + 1$ processes are needed to achieve f -fault tolerance. In the worst case, the f failing processes could accidentally generate the same reply. However, the remaining $f + 1$ will also produce the same answer, so the LPP just believes the majority.

In our design each map or reduce task scheduled by current LPP are assigned to independent WPGs. Let's try to explain this statement formally using map task as example (but is the same of any task). Let $i = 1, 2, 3 \dots n$, where $n \in \mathbb{Z}_+$. As MapReduce framework states, our system splits the input text file S in n parts $\{s_1, s_2, \dots, s_n\}$ called *split*; each split s_i is associated to a map task $m_i \in \{m_1, m_2, \dots, m_n\}$ every of which must be replicated on multiple WPs to make them arbitrary fault tolerant. Each map task m_i is assigned to a different WPG $g_i \in \{g_1, g_2, \dots, g_n\}$. As we will explain later, LPP, checking received outputs from each group, we can achieve arbitrary fault tolerance for every scheduled task. Obviously to tolerate at most f arbitrary failures of each task, each WPG must contains at least $2f + 1$ different processes. Then our design requires at least $n(2f + 1)$ different WPs.

This design is scalable in several ways: if the number of users or data input size grows, it should be possible to extend the system, at reasonable cost, adding more WPGs. If we want an higher level of arbitrary fault-tolerance, we should add more WP to each WPG, increasing the number PP backups too.

4 THE ALGORITHM

Given the system model described previously, a solution to tolerate at most f arbitrary faults during each task would be the following: current LPP assigns a given map/reduce task to a WPG contacting its $2f + 1$ WPs; then, after retrieving the outputs from replicas, picks the most voted results which is expected correct if no more than f arbitrary faults occur.

This simplistic solution is very expensive because it replicates everything $2f + 1$ times: task execution, communication of map task outputs and so on. We use a set of techniques to avoid these costs and do things efficiently:

Digest outputs

As explained above, in order to consider a map/reduce task correct, tolerating at most f faulty processes, $f + 1$ matching outputs, coming from at least $2f + 1$ WPs belonging to a given WPG, have to be received.

To validate task's output, since outputs computed by worker processes can be very large, in order to avoid pointless additional network traffic, we have adopted an approach according to which LPP fetches and compares *digests*, a bit string of a fixed size obtained using a hash function. The main advantage of this technique is that, digest size are many times smaller than task's data output, allowing us to reduce network traffic among servers, increasing system's performance.

Deferred execution

We believe that there is no point in always executing $2f + 1$ replicas for each task to usually obtain the same result.

To minimize both the number of copies of tasks executed and the overhead due to network traffic we have adopted a design called *deferred execution*: according to that solution current LPP starts only $f + 1$ replicas of the same task, checking if they all return the same result. If a time-out elapses, or some returned results do not match, more replicas (up to f) are started, until there are $f + 1$ matching replies.

We believe that this this design is also energy-saver and cost-effective since, supposing that arbitrary faults are uncommon, fewer WP are involved into computations respect to basic scheme and, in principle, is possible to start and use a fewer number of servers using auto-scaling techniques (at cost of some latency).

Data locality awareness

Since in order to achieve fault tolerance we have replicated all task $2f + 1$ times, data transfers among servers are replicated too, generating an huge amount of network traffic. Evidently, when large data set are involved, move data repeatedly among server became a bottleneck.

Therefore we have adopted a design capable to reduce

Since moving data repeatedly among server is bottleneck, especially when large data set are involved, in order to improve system performance, we have adopt a design aware of data locations and sizes in order to mitigate network traffic avoiding pointless data moving.

Our solution is based on a basic principles which states that "*moving computation towards data is cheaper than moving data towards computation*". When the map phase is fully done, all the network locations of the feeding nodes of every reducer will be known

Moving data repeatedly to distant nodes is becoming the bottleneck [23]. In this paper we rethink reduce task scheduling in Hadoop and suggest making Hadoop's reduce task scheduler aware of partitions' network locations and sizes in order to mitigate network traffic.

A key question is how to schedule reduce tasks at Task Trackers so as to diminish shuffled data and improve MapReduce performance. One of Hadoop's basic principles is: moving computation towards data is cheaper than moving data towards computation. Such a principle is employed by Hadoop when scheduling map tasks but bypassed when scheduling reduce tasks. MapReduce is aware of the network locations of splits (inputstomappers)andleveragessuchinformationtoschedule mappers nearby splits. In contrast, MapReduce is oblivious to the network locations of partitions (inputs to reducers) and does not schedule reducers nearby partitions. Thus, similar to map task scheduling, we suggest making MapReduce aware of partitions network locations in order to apply locality to reduce task scheduling

Crash failure detection

To manage and detect workers processes crash faults we have use some features of Apache Zookeeper.

As known, Apache ZooKeeper has the notion of *ephemeral nodes*, that is special znodes which exists as long as the *session* that created the znode is active. When session expiration occurs, Zookeeper cluster will delete all ephemeral nodes owned by that session and immediately notify all connected clients of the change.

When any client establishes a session with a Zookeeper cluster, a time-out value is used by the cluster to determine when the client's session expires. Expirations usually happens when the cluster does not hear from the client within the specified session time-out period (i.e. no heartbeat).

Notice that session is kept alive by requests sent by client, therefore is critical that client will send PING request to keep the session alive. This PING request not only allows the ZooKeeper server to know that the client is still active, but it also allows the client to verify that its connection to the ZooKeeper server is still active.

Using this design, is very easy to keep check the status of worker processes.

REFERENCES

- [1] Pedro Costa, Marcelo Pasin, Alysson N. Bessani, and Miguel Correia. 2011. Byzantine Fault-Tolerant MapReduce: Faults Are Not Just Crashes. (2011).
- [2] Mohammad Hammoud and Majd F. Sakr. 2011. Locality-Aware Reduce Task Scheduling for MapReduce. (2011).
- [3] IBM. [n.d.]. *What is MapReduce?* Retrieved September 2, 2019 from <https://www.ibm.com/analytics/hadoop/mapreduce>
- [4] E. B. Nightingale, J. R. Douceur, and V. Orgovan. 2011. Cycles, cells and platters: an empirical analysis of hardware failures on a million consumer PCs. *Proceedings of the EuroSys 2011 Conference* (2011), 343—356.
- [5] B. Schroeder and G. A. Gibson. 2007. Understanding failures in petascale computers. *Journal of Physics: Conference Series* (2007), 78.
- [6] B. Schroeder, E. Pinheiro, and W.-D. Weber. 2009. DRAM errors in the wild: a large-scale field study. *In Proceedings of the 11th International Joint Conference on Measurement and Modeling of Computer Systems* (2009), 193—204.
- [7] Maarten van Steen and Andrew S. Tanenbaum. Version 3.01 (2017). *Distributed Systems*.