

Arbitrary Fault-Tolerant and Locality-Aware MapReduce

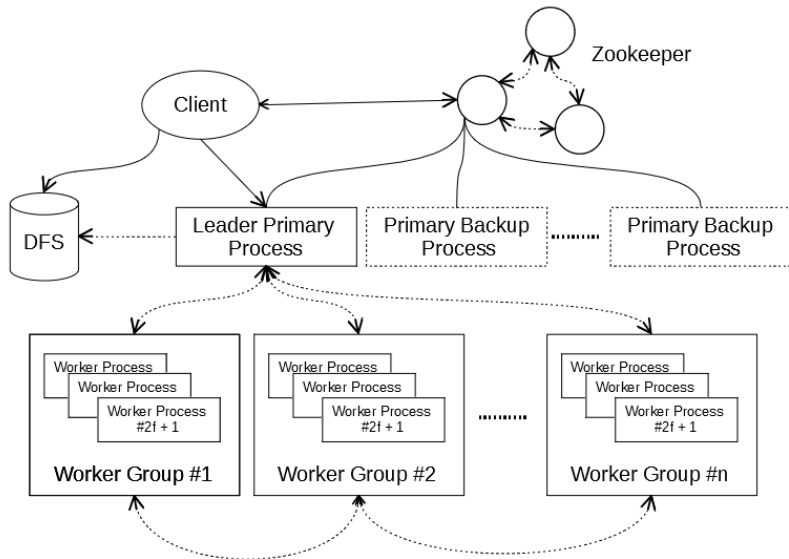
Sistemi Distribuiti e Cloud Computing 2018-2019

Andrea Graziani - 0273395

Università degli Studi di Roma "Tor Vergata"
FACOLTA' DI INGEGNERIA
Corso di Laurea Magistrale in Ingegneria Informatica

2 ottobre 2019

System Architecture



System Architecture

- What assumptions we have made?
- We distinguish three process types:
 - *Client Process*
 - *Primary Process* (PP)
 - *Worker Process* (WP)
- How we made the system fault tolerant? What kind of fault tolerance can be achieved with this design?
- What kind of semantics do we have adopted?
- Is our system scalable?
- Our system can be elastic? **Yes... but**
- What are the advantages/disadvantages of our solution?

System features

Deferred Execution

- By default our system is design to use the *Deferred Execution* approach during task scheduling.
 - What's meant by deferred execution and how it works?
 - What are the advantages of this approach? And the disadvantages?
 - What happens if LPP/WPs crashes?
 - Why *Locality-Aware*?
 - Why *State-Aware*?

System features

Speculative Execution

- In order to improve system's performance we have adopt an approach based on the so called *Speculative Execution*
 - What's meant by speculative execution? How it works?
 - What happens if the input used during Reduce-Phase is detected as incorrect?
 - What happens if LPP/WPs crashes?
- Can we do better? **Yes...but...**

System features

Data-Locality-Aware Reduce Scheduling and Shuffle

- Our system is designed to schedule reduce task considering data locality in order to reduce the overall amount of transferred data between WPs: this approach is called *Data-Locality-Aware Reduce Scheduling*.
 - "Moving computation towards data is cheaper than moving data towards computation".
 - How it works?

System features

Other features

- Digest output.
- Leader election and crash failure detection based on Apache Zookeeper.

Grazie per l'attenzione!